

clustering-Text-K-MEANS-NLP

March 23, 2024

```
[1]: from sklearn.datasets import fetch_20newsgroups
from nltk.tokenize import word_tokenize #Used to extract words from documents
from nltk.stem import WordNetLemmatizer #Used to lemmatize words
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import Normalizer
from sklearn import metrics

from sklearn.cluster import KMeans

import sys
from time import time

import pandas as pd
import numpy as np
```

```
[2]: # Selected 3 categories from the 20 newsgroups dataset

categories = [
    'talk.religion.misc',
    'comp.graphics',
    'sci.space',
]

print("Loading 20 newsgroups dataset for categories:")
print(categories)
```

Loading 20 newsgroups dataset for categories:
['talk.religion.misc', 'comp.graphics', 'sci.space']

```
[3]: df = fetch_20newsgroups(subset='all', categories=categories,
                             shuffle=False, remove=('headers', 'footers',
↪ 'quotes'))
```

```
[4]: labels = df.target
true_k = len(np.unique(labels)) ## This should be 3 in this example
print(true_k)
```

```
[5]: #Perform Lemmatization
```

```
[6]: lemmatizer = WordNetLemmatizer()
for i in range(len(df.data)):
    word_list = word_tokenize(df.data[i])
    lemmatized_doc = ""
    for word in word_list:
        lemmatized_doc = lemmatized_doc + " " + lemmatizer.lemmatize(word)
    df.data[i] = lemmatized_doc
```

```
[7]: print(df.data[1])
```

In regard to fractal commpression , I have seen 2 fractal compressed `` movie '' . They were both fairly impressive . The first one wa a 64 gray scale `` movie '' of Casablanca , it wa 1.3MB and had 11 minute of 13 fps video . It wa a little grainy but not bad at all . The second one I saw wa only 3 minute but it had 8 bit color with 10fps and measured in at 1.2MB . I consider the fractal movie a practical thing to explore . But unlike many other format out there , you do end up losing resolution . I do n't know what kind of software/hardware wa used for creating the `` movie '' I saw but the guy that showed them to me said it took 5-15 minute per frame to generate . But a I said above playback wa 10 or more frame per second . And how else could you put 11 minute on one floppy disk ?

```
[8]: #We next convert our corpus into tf-idf vectors. We remove common stop words,
↳terms with very low document frequency (many of them are numbers
#or misspells), accents.
```

```
[9]: vectorizer = TfidfVectorizer(strip_accents='unicode', stop_words='english',
↳min_df=2) ## Corpus is in English
X = vectorizer.fit_transform(df.data)
```

```
[10]: print(X.shape)
```

(2588, 14439)

```
[11]: #Clustering using standard k-means¶
```

```
[12]: #We choose a value that reflects our knowledge about the data
```

```
[13]: #We may try several value, possibly in increasing order
```

```
[14]: #we set k=3
```

```
[15]: km = KMeans(n_clusters=true_k, init='k-means++', max_iter=100)
t0 = time()
km.fit(X)
print("done in %0.3fs" % (time() - t0))
```

done in 0.162s

```
/home/nmit/anaconda3/lib/python3.11/site-  
packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of  
'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init'  
explicitly to suppress the warning  
warnings.warn(
```

```
[16]: #Standard measures of cluster quality
```

```
[17]: print("Homogeneity: %0.3f" % metrics.homogeneity_score(labels, km.labels_))  
print("Completeness: %0.3f" % metrics.completeness_score(labels, km.labels_))  
print("V-measure: %0.3f" % metrics.v_measure_score(labels, km.labels_))  
print("Adjusted Rand-Index: %.3f"  
      % metrics.adjusted_rand_score(labels, km.labels_))  
print("Silhouette Coefficient: %0.3f"  
      % metrics.silhouette_score(X, km.labels_, sample_size=1000))
```

```
Homogeneity: 0.293  
Completeness: 0.401  
V-measure: 0.339  
Adjusted Rand-Index: 0.210  
Silhouette Coefficient: 0.009
```

```
[18]: #Identify the 10 most relevant terms in each cluster
```

```
[20]: centroids = km.cluster_centers_.argsort()[:, :-1] ## Indices of largest  
      ↪centroids' entries in descending order  
#vectorizer.get_feature_names use instead vectorizer.get_feature_names_out  
terms = vectorizer.get_feature_names_out()  
for i in range(true_k):  
    print("Cluster %d:" % i, end='')  
    for ind in centroids[i, :10]:  
        print(' %s' % terms[ind], end='')  
    print()
```

```
Cluster 0: god jesus christian bible wa people christians say christ did  
Cluster 1: wa space just like think ha time know did people  
Cluster 2: file image thanks format program know graphic ftp bit gif
```

```
[21]: #Visualization
```

```
[22]: from wordcloud import WordCloud  
import matplotlib.pyplot as plt
```

```
[23]: def frequencies_dict(cluster_index):  
    if cluster_index > true_k - 1:  
        return  
    term_frequencies = km.cluster_centers_[cluster_index]
```

```
sorted_terms = centroids[cluster_index]
frequencies = {terms[i]: term_frequencies[i] for i in sorted_terms}
return frequencies
```

```
[24]: def makeImage(frequencies):

    wc = WordCloud(background_color="white", max_words=50)
    # generate word cloud
    wc.generate_from_frequencies(frequencies)

    # show
    plt.imshow(wc, interpolation="bilinear")
    plt.axis("off")
    plt.show()
```

```
[25]: for i in range(true_k):
    freq = frequencies_dict(i)
    makeImage(freq)
    print()
```



